

```
# Loading relevant libraries
```

```
library(readxl)
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2     3.5.1      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr       1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(TTR)
```

```
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##   method      from
```

```
## as.zoo.data.frame zoo
```

```
# Importing the death data from excel
```

```
uk_death <- read_excel('Vital statistics in the UK.xlsx', sheet = 3, skip = 5)
```

```
# Checking the head of the data frame
```

```
head(uk_death)
```

```
## # A tibble: 6 x 7
```

```
##   Year Number of deaths: United~1 Number of deaths: En~2 Number of deaths: En~3
```

```
##   <dbl> <chr>                                <dbl> <chr>
```

```
## 1  2021 667479                                586334 549349
```

```
## 2  2020 689629                                607922 569700
```

```
## 3  2019 604707                                530841 496370
```

```
## 4  2018 616014                                541589 505859
```

```
## 5  2017 607172                                533253 498882
```

```
## 6  2016 597206                                525048 490791
```

```
## # i abbreviated names: 1: 'Number of deaths: United Kingdom',
```

```
## #   2: 'Number of deaths: England and Wales', 3: 'Number of deaths: England'
```

```
## # i 3 more variables: 'Number of deaths: Wales' <chr>,
```

```
## #   'Number of deaths : Scotland' <chr>,
```

```
## #   'Number of deaths: Northern Ireland' <chr>
```

```
# Selecting columns needed (Year, Number of deaths: United Kingdom)
```

```
death_uk <- uk_death %>%
```

```
  select(Year, `Number of deaths: United Kingdom`)
```

```
# Checking the first 6 entries.
```

```
head(death_uk)
```

```
## # A tibble: 6 x 2
##   Year 'Number of deaths: United Kingdom'
##   <dbl> <chr>
## 1  2021 667479
## 2  2020 689629
## 3  2019 604707
## 4  2018 616014
## 5  2017 607172
## 6  2016 597206
```

```
# Checking the last 6 entries.
tail(death_uk)
```

```
## # A tibble: 6 x 2
##   Year 'Number of deaths: United Kingdom'
##   <dbl> <chr>
## 1  1843 :
## 2  1842 :
## 3  1841 :
## 4  1840 :
## 5  1839 :
## 6  1838 :
```

Notice that some observations are ‘:’

```
# Checking the structure of the data frame
str(death_uk)
```

```
## tibble [184 x 2] (S3: tbl_df/tbl/data.frame)
##  $ Year                : num [1:184] 2021 2020 2019 2018 2017 ...
##  $ Number of deaths: United Kingdom: chr [1:184] "667479" "689629" "604707" "616014" ...
```

from the above, the ‘Number of deaths: United Kingdom’ column is stored as ‘chr’

```
# Cleaning and preparing the data for time series analysis.
death_uk <- death_uk %>%
  rename(no_of_deaths = `Number of deaths: United Kingdom`) %>%
  filter(no_of_deaths != ':') %>%
  arrange(Year) %>%
  select(no_of_deaths)
```

```
# Converting the data to a time series
death_uk$no_of_deaths <- as.integer(death_uk$no_of_deaths)
death_uk_ts = ts(death_uk, frequency = 1, start = 1887)
death_uk_ts
```

```
## Time Series:
## Start = 1887
## End = 2021
## Frequency = 1
##      no_of_deaths
```


##	[1,]	629287
##	[2,]	605899
##	[3,]	615033
##	[4,]	665758
##	[5,]	696490
##	[6,]	661273
##	[7,]	673722
##	[8,]	593808
##	[9,]	676110
##	[10,]	620108
##	[11,]	645630
##	[12,]	654812
##	[13,]	685510
##	[14,]	695867
##	[15,]	655646
##	[16,]	636650
##	[17,]	613726
##	[18,]	651301
##	[19,]	617516
##	[20,]	629955
##	[21,]	625271
##	[22,]	621427
##	[23,]	614910
##	[24,]	578091
##	[25,]	620868
##	[26,]	580977
##	[27,]	600554
##	[28,]	611970
##	[29,]	666322
##	[30,]	599621
##	[31,]	589416
##	[32,]	715246
##	[33,]	602188
##	[34,]	555326
##	[35,]	544140
##	[36,]	579480
##	[37,]	526858
##	[38,]	563891
##	[39,]	558132
##	[40,]	536411
##	[41,]	568655
##	[42,]	543664
##	[43,]	623231
##	[44,]	536860
##	[45,]	573908
##	[46,]	567986
##	[47,]	579467
##	[48,]	558072
##	[49,]	561324
##	[50,]	580942
##	[51,]	597798
##	[52,]	559598
##	[53,]	581857
##	[54,]	673253

##	[55,]	627378
##	[56,]	562356
##	[57,]	585582
##	[58,]	573570
##	[59,]	567027
##	[60,]	573361
##	[61,]	600728
##	[62,]	546002
##	[63,]	589876
##	[64,]	590136
##	[65,]	632786
##	[66,]	573806
##	[67,]	577220
##	[68,]	578400
##	[69,]	595916
##	[70,]	597981
##	[71,]	591200
##	[72,]	604040
##	[73,]	606115
##	[74,]	603328
##	[75,]	631788
##	[76,]	636051
##	[77,]	654288
##	[78,]	611130
##	[79,]	627798
##	[80,]	643754
##	[81,]	616710
##	[82,]	655998
##	[83,]	659537
##	[84,]	655385
##	[85,]	645078
##	[86,]	673938
##	[87,]	669692
##	[88,]	667359
##	[89,]	662477
##	[90,]	680799
##	[91,]	655143
##	[92,]	667177
##	[93,]	675576
##	[94,]	661519
##	[95,]	657974
##	[96,]	662801
##	[97,]	659101
##	[98,]	644918
##	[99,]	670656
##	[100,]	660735
##	[101,]	644342
##	[102,]	649178
##	[103,]	657733
##	[104,]	641799
##	[105,]	646181
##	[106,]	634238
##	[107,]	658194
##	[108,]	626222

```
## [109,]      641712
## [110,]      638879
## [111,]      632517
## [112,]      627592
## [113,]      629476
## [114,]      610579
## [115,]      604393
## [116,]      608045
## [117,]      612085
## [118,]      584791
## [119,]      582964
## [120,]      572224
## [121,]      574687
## [122,]      579697
## [123,]      559617
## [124,]      561666
## [125,]      552232
## [126,]      569024
## [127,]      576458
## [128,]      570341
## [129,]      602782
## [130,]      597206
## [131,]      607172
## [132,]      616014
## [133,]      604707
## [134,]      689629
## [135,]      667479
```

```
# Plotting the initial number of deaths from year 1887 - 2021
plot.ts(death_uk_ts, main='Time series of number of deaths in UK (1887 -2021)')
```


Time series of number of deaths in UK (1887 –2021)



The time series appears non seasonal and can probably be described using an additive model. Time series is non seasonal, but has trend and irregular components.

```
# Estimating trend by smoothing using a simple moving average of order 10.
plot.ts(SMA(death_uk_ts, n=10),
        main='Time series showing trend of number of deaths in UK',
        ylab = 'no of deaths (SMA)')
```

Time series showing trend of number of deaths in UK



TIME SERIES MODELLING


MODEL 1 -FORECASTING USING SMOOTHING The time series can be described by an additive model, it has trend with no seasonality, therefore: We can use Holt's Exponential Smoothing.

```
# Fitting a predictive model using Holt-Winters
death_uk_ts_forecast <- HoltWinters(death_uk_ts, gamma = FALSE)
death_uk_ts_forecast
```

```
## Holt-Winters exponential smoothing with trend and without seasonal component.
##
## Call:
## HoltWinters(x = death_uk_ts, gamma = FALSE)
##
## Smoothing parameters:
##  alpha: 0.4796249
##  beta : 0.1534181
##  gamma: FALSE
##
## Coefficients:
##           [,1]
## a 662949.045
## b   9829.915
```


An alpha value approximately 0.48, is just right in the middle of 0 and 1, which means that 48% of the weight is given to the most recent observation when estimating the level. A beta value of 0.15 means more weight (85%) is given to the previous trend estimate (not the most recent).

```
# plotting both observed and fitted data from HoltWinters forecast.  
plot(death_uk_ts_forecast)
```



```
# forecasting the next 10 years.  
death_uk_ts_forecast2 <- forecast(death_uk_ts_forecast, h=10)  
plot(death_uk_ts_forecast2)
```


Forecasts from HoltWinters



The forecast in blue. The purple area is the 80% prediction interval The grey area is the 95% prediction interval

```
# forecasted data with the 80% and 85% intervals.
death_uk_ts_forecast2
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 2022	672779.0	633216.8	712341.1	612273.9	733284.0
## 2023	682608.9	637396.5	727821.3	613462.5	751755.3
## 2024	692438.8	640872.7	744004.9	613575.2	771302.3
## 2025	702268.7	643729.7	760807.7	612741.0	791796.4
## 2026	712098.6	646035.3	778161.9	611063.5	813133.7
## 2027	721928.5	647843.1	796014.0	608624.6	835232.5
## 2028	731758.4	649195.5	814321.4	605489.3	858027.6
## 2029	741588.4	650126.4	833050.3	601709.4	881467.3
## 2030	751418.3	650663.6	852173.0	597327.2	905509.3
## 2031	761248.2	650829.5	871666.9	592377.4	930119.0

```
# Sum of square error
death_uk_ts_forecast2$SSE
```

```
## [1] 127326590931
```


```
# Root Mean Square Error for Holt-Winters
RMSE_HW = sqrt(mean(death_uk_ts_forecast2$residuals^2, na.rm = TRUE))
RMSE_HW
```

```
## [1] 30940.96
```

```
# Mean absolute percentage error MAPE for Holt-Winters
MAPE_HW = mean((abs(death_uk_ts_forecast2$residuals/death_uk_ts)*100), na.rm=TRUE)
MAPE_HW
```

```
## [1] 3.537774
```

```
# ACF and Ljung box test
acf(death_uk_ts_forecast2$residuals, lag.max=20 , na.action = na.pass)
```




```
Box.test(death_uk_ts_forecast2$residuals, lag=20, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data: death_uk_ts_forecast2$residuals
## X-squared = 31.015, df = 20, p-value = 0.055
```

The P-value for the Ljung-test is 0.055, there is little evidence of non-zero auto correlations in the in-sample forecast errors at lags 1-20.

```
# plotting the forecast errors to check for constant variance
plot.ts(death_uk_ts_forecast2$residuals, main='Holt-Winters Forecast Errors',
        ylab = 'Residuals')
```



```
# function to plot forecast errors and overlay normal distributed data
plotForecastErrors <- function(forecasterrors)
{
  # make a histogram of the forecast errors:
  mybinsize <- IQR(forecasterrors)/4
  mysd <- sd(forecasterrors)
  mymin <- min(forecasterrors) - mysd*5
  mymax <- max(forecasterrors) + mysd*3
  # generate normally distributed data with mean 0 and standard deviation mysd
  mynorm <- rnorm(10000, mean=0, sd=mysd)
  mymin2 <- min(mynorm)
  mymax2 <- max(mynorm)
  if (mymin2 < mymin) { mymin <- mymin2 }
  if (mymax2 > mymax) { mymax <- mymax2 }
  # make a red histogram of the forecast errors, with the normally distributed data overlaid:
  mybins <- seq(mymin, mymax, mybinsize)
  hist(forecasterrors, col="red", freq=FALSE, breaks=mybins)
  # freq=FALSE ensures the area under the histogram = 1
  # generate normally distributed data with mean 0 and standard deviation mysd
```

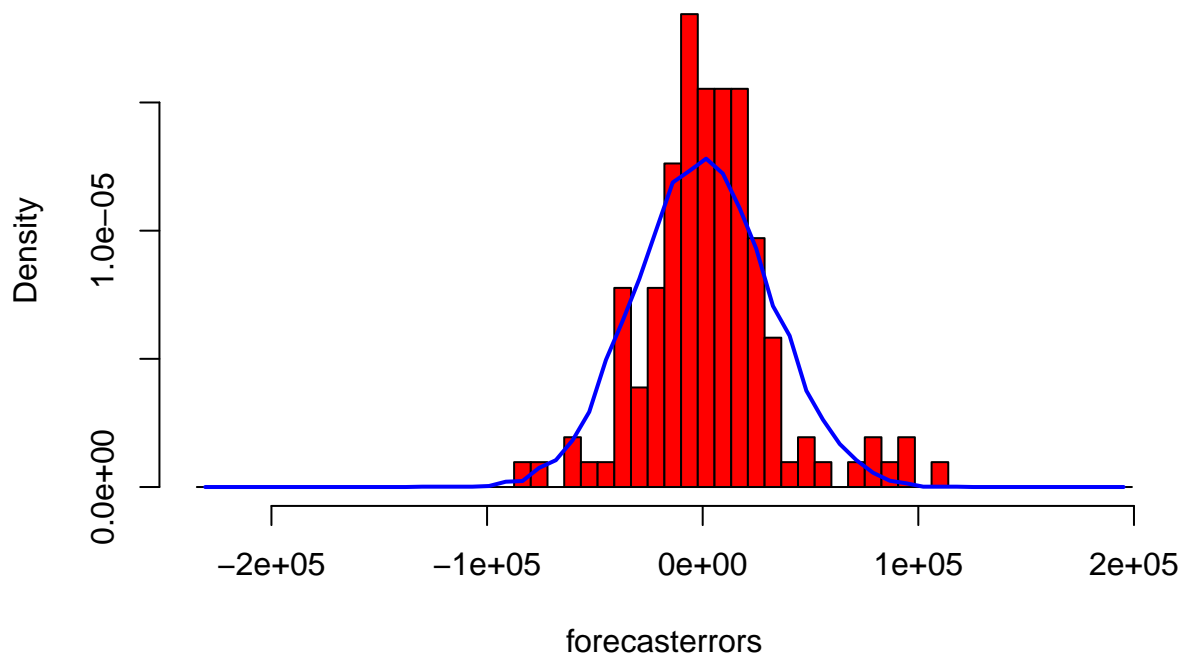
```
myhist <- hist(mynorm, plot=FALSE, breaks=mybins)
# plot the normal curve as a blue line on top of the histogram of forecast errors:
points(myhist$mids, myhist$density, type="l", col="blue", lwd=2)
}
```

```
# removing NA values from the residuals
```

```
death_uk_ts_forecast2$residuals <- death_uk_ts_forecast2$residuals[!is.na(death_uk_ts_forecast2$residuals)]
```

```
# plotting if the forecast errors to check if normally distributed with mean zero
plotForecastErrors(death_uk_ts_forecast2$residuals)
```

Histogram of forecasterrors



```
# library to import Augmented Dickey-Fuller Test
library(tseries)
```

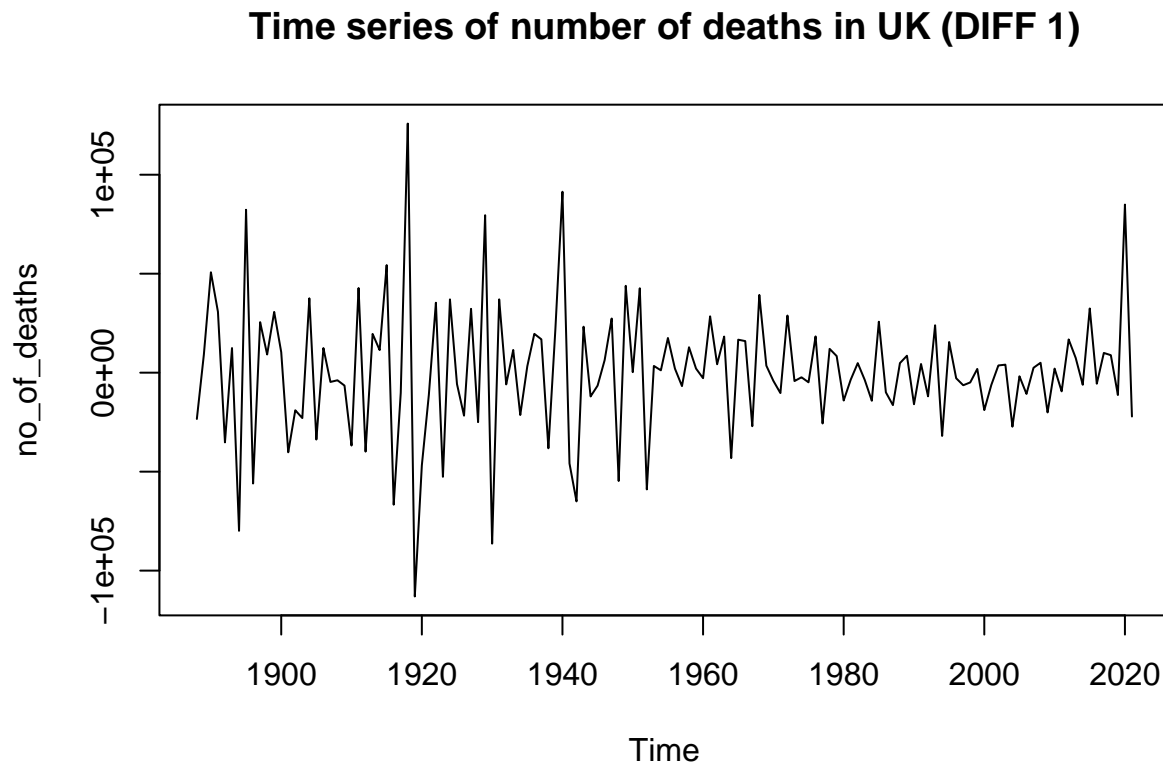
```
# Augmented Dickey-Fuller Test
adf.test(death_uk_ts)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: death_uk_ts
## Dickey-Fuller = -2.3315, Lag order = 5, p-value = 0.4386
## alternative hypothesis: stationary
```

Test if series is stationary, P-value is greater than 0.05, therefore we fail to reject null hypothesis.

```
# Differencing the time series to make it stationary
death_uk_ts_diff1 <- diff(death_uk_ts, differences = 1)

# Plotting the series with difference 1.
plot(death_uk_ts_diff1, main='Time series of number of deaths in UK (DIFF 1)')
```



The plot appears stationary in mean

```
# Augmented Dickey-Fuller Test for difference 1
adf.test(death_uk_ts_diff1)
```


```
## Warning in adf.test(death_uk_ts_diff1): p-value smaller than printed p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: death_uk_ts_diff1
## Dickey-Fuller = -6.1896, Lag order = 5, p-value = 0.01
## alternative hypothesis: stationary
```

Test if series is stationary, P-value is less than 0.05, therefore we reject null hypothesis. Difference 1 is stationary.

SELECTING ARIMA MODEL.

```
# Plotting the correlogram for diff1
acf(death_uk_ts_diff1, lag.max = 20)
```




```
acf(death_uk_ts_diff1, lag.max = 20, plot = FALSE)
```

```
##
## Autocorrelations of series 'death_uk_ts_diff1', by lag
##
##      0      1      2      3      4      5      6      7      8      9     10
## 1.000 -0.446 -0.004 -0.042  0.177 -0.120 -0.116  0.244 -0.180  0.086 -0.102
##     11     12     13     14     15     16     17     18     19     20
## 0.231 -0.166 -0.068  0.147 -0.090  0.072 -0.202  0.260 -0.114 -0.014
```

from the correlogram, the autocorrelation at lag 1 (-0.446) exceeds the significance bounds. so a Moving average model of order 1 - ARMA(0,1) can be used which is also a ARIMA(0,1,1) with difference 1.

```
# Plotting the partial correlogram for diff1
pacf(death_uk_ts_diff1, lag.max = 20)
```

Series death_uk_ts_diff1



```
pacf(death_uk_ts_diff1, lag.max = 20, plot = FALSE)
```

```
##
## Partial autocorrelations of series 'death_uk_ts_diff1', by lag
##
##      1      2      3      4      5      6      7      8      9     10     11
## -0.446 -0.253 -0.208  0.079  0.008 -0.182  0.128 -0.079  0.029 -0.044  0.143
##      12     13     14     15     16     17     18     19     20
##  0.054 -0.112  0.038 -0.062  0.050 -0.133  0.017  0.067 -0.028
```

The partial autocorrelation at lags 1,2,and 3 exceeds the significance bounds. an Auto regressive model of order 3 is possible. ARIMA(3,1,0)

From the principle of parsimony (fewer is better).

MODEL 2 - MOVING AVERAGE MODEL OF ORDER 1 - ARIMA(0,1,1)

```
# Moving average model of order 1 and difference 1.
death_uk_ts_ma <- arima(death_uk_ts, order = c(0,1,1))
death_uk_ts_ma
```


```
##
## Call:
## arima(x = death_uk_ts, order = c(0, 1, 1))
##
```

```
## Coefficients:
##      ma1
##      -0.5913
## s.e.    0.0695
##
## sigma^2 estimated as 812067647:  log likelihood = -1564.86,  aic = 3133.73
```

```
# forecasting the next 10 years using moving average.
death_uk_ts_ma_forecast <- forecast(death_uk_ts_ma, h =10)
death_uk_ts_ma_forecast
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 2022	650827	614306.9	687347.1	594974.3	706679.7
## 2023	650827	611374.2	690279.8	590489.2	711164.8
## 2024	650827	608645.0	693009.1	586315.1	715338.9
## 2025	650827	606081.9	695572.2	582395.2	719258.8
## 2026	650827	603657.8	697996.2	578688.0	722966.1
## 2027	650827	601352.4	700301.6	575162.2	726491.9
## 2028	650827	599149.8	702504.3	571793.5	729860.6
## 2029	650827	597037.2	704616.8	568562.6	733091.4
## 2030	650827	595004.6	706649.5	565454.0	736200.1
## 2031	650827	593043.4	708610.6	562454.6	739199.5

```
# 10 year forecast plot
plot(death_uk_ts_ma_forecast)
```




```
# Evaluation for ARIMA(0,1,1)
AIC(death_uk_ts_ma)
```

```
## [1] 3133.728
```

```
BIC(death_uk_ts_ma)
```

```
## [1] 3139.524
```


```
accuracy(death_uk_ts_ma)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 456.7854 28391.11 20199.13 -0.09581522 3.280053 0.858626
##           ACF1
## Training set -0.01616067
```

```
# ACF and Ljung box test
```

```
acf(death_uk_ts_ma_forecast$residuals, lag.max=20, na.action = na.pass)
```

Series death_uk_ts_ma_forecast\$residuals




```
Box.test(death_uk_ts_ma_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##
```

```
## Box-Ljung test
##
## data: death_uk_ts_ma_forecast$residuals
## X-squared = 31.243, df = 20, p-value = 0.05206
```


The P-value for the Ljung-test is 0.052, there is little evidence of non-zero auto correlations in the forecast errors at lags 1-20.

```
# plotting the forecast errors to check for constant variance
plot.ts(death_uk_ts_ma_forecast$residuals, main = 'ARIMA(0,1,1) Forecast Errors',
        ylab = 'Residuals')
```



Appears to have mean 0 and constant Variance

```
# plotting if the forecast errors to check if normally distributed with mean zero
plotForecastErrors(death_uk_ts_ma_forecast$residuals)
```



Appears normally distributed with mean 0

MODEL 3 - ARIMA(3,1,0)

```
# ARIMA model (3,1,0)
death_uk_ts_ar <- arima(death_uk_ts, order = c(3,1,0))
death_uk_ts_ar
```


```
##
## Call:
## arima(x = death_uk_ts, order = c(3, 1, 0))
##
## Coefficients:
##          ar1      ar2      ar3
##      -0.6108  -0.3798  -0.2166
## s.e.   0.0844   0.0968   0.0871
##
## sigma^2 estimated as 794265312:  log likelihood = -1563.41,  aic = 3134.83
```

```
# forecasting the next 10 years.
death_uk_ts_ar_forecast <- forecast(death_uk_ts_ar, h =10)
death_uk_ts_ar_forecast
```

```
##      Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## 2022      651205.7 615088.1 687323.3 595968.6 706442.8
## 2023      651163.3 612406.4 689920.1 591889.8 710436.8
```

```
## 2024      662167.0 621021.6 703312.5 599240.6 725093.5
## 2025      658987.0 615355.4 702618.6 592258.2 725715.7
## 2026      656759.6 609186.9 704332.2 584003.5 729515.6
## 2027      656944.3 606743.0 707145.7 580168.0 733720.7
## 2028      658366.2 605682.6 711049.8 577793.6 738938.8
## 2029      657910.0 602803.2 713016.9 573631.4 742188.7
## 2030      657608.6 600025.6 715191.7 569543.0 745674.3
## 2031      657658.0 597808.2 717507.8 566125.6 749190.4
```

```
# 10 year forecast plot
plot(death_uk_ts_ar_forecast)
```



```
# Evaluation for ARIMA(3,1,0)
AIC(death_uk_ts_ar)
```

```
## [1] 3134.827
```

```
BIC(death_uk_ts_ar)
```


```
## [1] 3146.418
```

```
accuracy(death_uk_ts_ar)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 535.029 28078.19 20038.88 -0.07547298 3.256082 0.8518145
##           ACF1
## Training set 0.01360795
```

```
# ACF and Ljung box test
acf(death_uk_ts_ar_forecast$residuals, lag.max=20 , na.action = na.pass)
```

Series death_uk_ts_ar_forecast\$residuals




```
Box.test(death_uk_ts_ar_forecast$residuals, lag=20, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data: death_uk_ts_ar_forecast$residuals
## X-squared = 24.19, df = 20, p-value = 0.2342
```

The P-value for the Ljung-test is 0.057, there is little evidence of non-zero auto correlations in the forecast errors at lags 1-20.

```
# plotting the forecast errors to check for constant variance
plot.ts(death_uk_ts_ar_forecast$residuals, main = 'ARIMA(3,1,0) Forecast Errors',
        ylab = 'Residuals')
```

ARIMA(3,1,0) Forecast Errors



```
# plotting if the forecast errors to check if normally distributed with mean zero  
plotForecastErrors(death_uk_ts_ar_forecast$residuals)
```

Histogram of forecasterrors

